

# **Alertmanager** **on its way to high availability**

**Frederic Branczyk**

Software Engineer at CoreOS

Prometheus/Alertmanager/Kubernetes

@brancz

# What that means?

- From alert to notification
- High availability contract
- High availability implementation
- Implications on operating HA Alertmanager

# Where does CoreOS fit in?

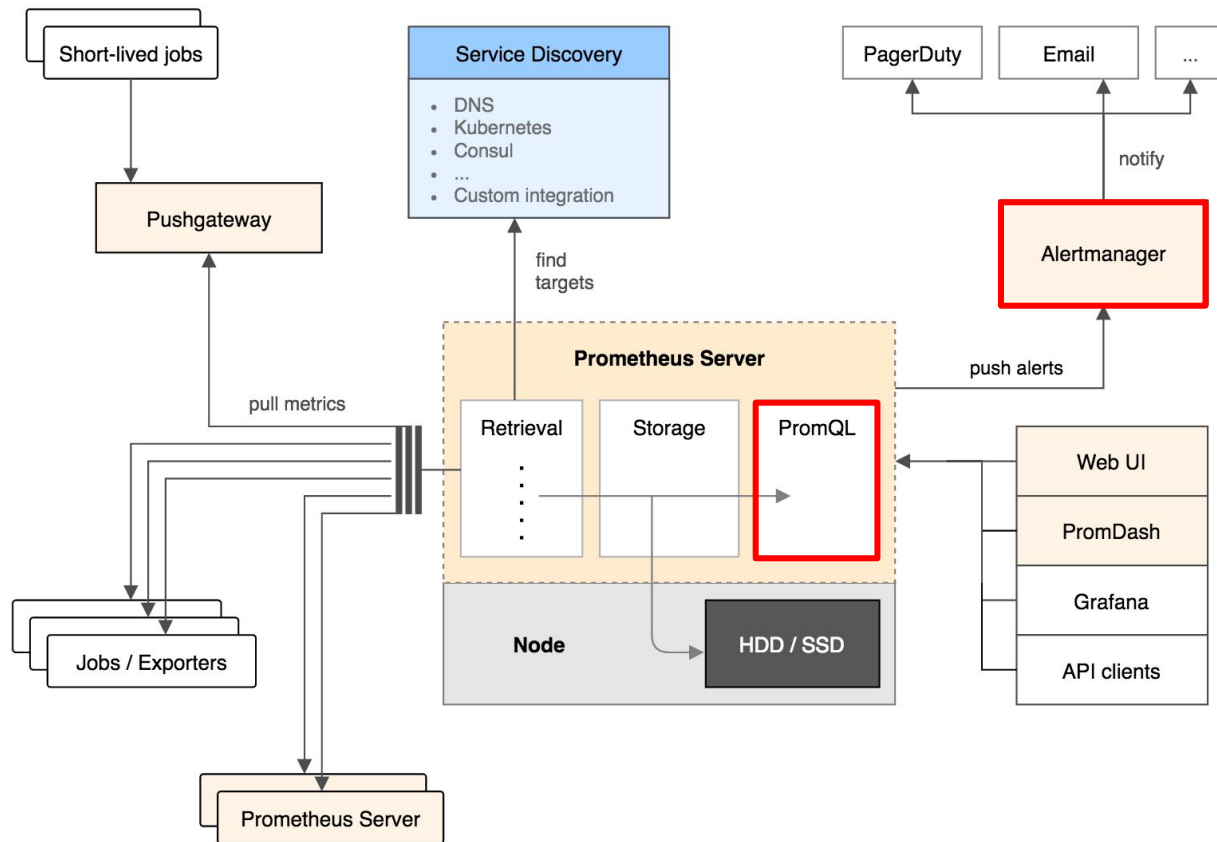
- Automating Monitoring infrastructure
- Prometheus + Kubernetes



# Alertmanager Features

- **Receives and groups alerts**
- **Deduplicates alerts**
- **Sends notifications to providers**
  - **Pagerduty, email, Slack, etc.**
- **Silencing**

# Prometheus & Alertmanager



Alerting Rule



Alerting Rule



...



Alerting Rule



Alerting Rule



```
04:11 hey, HighLatency, service="X", zone="eu-west", path=/user/profile, method=GET
04:11 hey, HighLatency, service="X", zone="eu-west", path=/user/settings, method=GET
04:11 hey, HighLatency, service="X", zone="eu-west", path=/user/settings, method=GET
04:11 hey, HighErrorRate, service="X", zone="eu-west", path=/user/settings, method=POST
04:12 hey, HighErrorRate, service="X", zone="eu-west", path=/user/profile, method=GET
04:13 hey, HighLatency, service="X", zone="eu-west", path=/index, method=POST
04:13 hey, CacheServerSlow, service="X", zone="eu-west", path=/user/profile, method=POST
. . .
04:15 hey, HighErrorRate, service="X", zone="eu-west", path=/comments, method=GET
04:15 hey, HighErrorRate, service="X", zone="eu-west", path=/user/profile, method=POST
```

# Grouped in one notification

- 3 x HighLatency
- 10 x HighErrorRate
- 2 x CacheServerSlow
- (+individual Alerts)

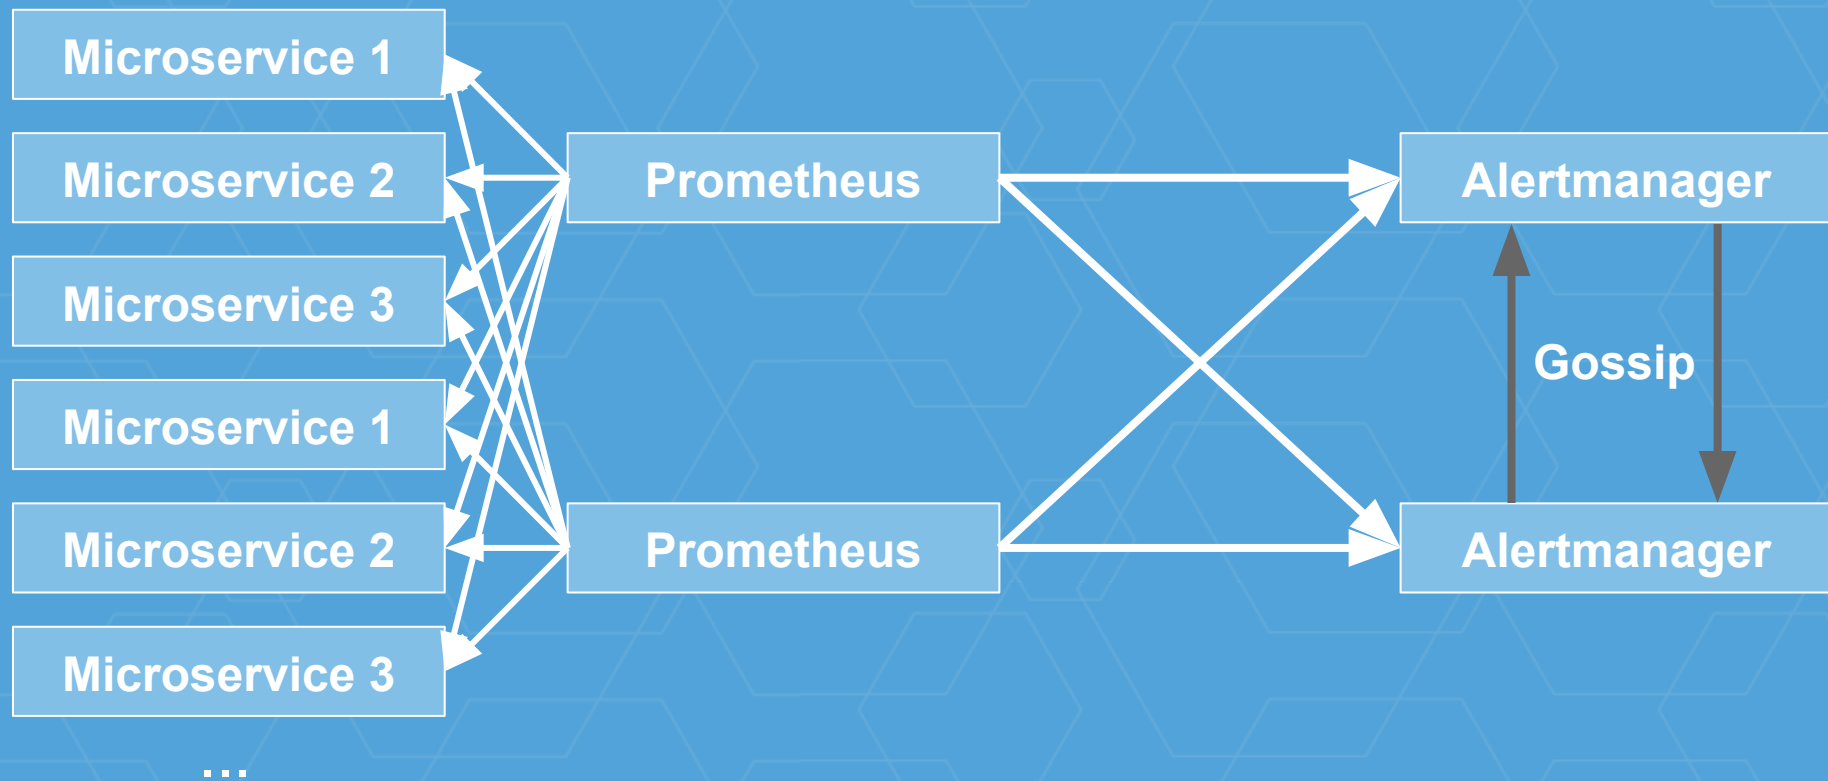
# Alertmanager

Reliably sends  
notifications



# High Availability

# Infrastructure Scaling Story



# Why decoupled?

- **Keep Prometheus alerting simple**
- **High availability of Prometheus**
- **No state sharing between Prometheus**

# Example Alerting Rule

```
ALERT NoLeader
IF etcd_has_leader == 0
FOR 10m
LABELS {
    severity = "warning"
}
ANNOTATIONS {
    summary = "etcd no leader",
    description = "etcd instance has no leader",
}
```

# Alert Evaluation in Prometheus



Rule 1      • Evaluate Rule/Alert

Rule 2      • Fire alert against Alertmanager

Rule 3

...

Repeat in *\*rule evaluation interval\**

# Simple configuration

```
global:
  resolve_timeout: 5m

route:
  group_by: ['job']
  group_wait: 10s
  group_interval: 10s
  repeat_interval: 1h
  receiver: 'webhook'
receivers:
- name: 'webhook'
  webhook_configs:
  - url: 'http://127.0.0.1:5001/'
```

- Resolve alerts in 5m
  - Group by *job* label
  - Group for 10 seconds
  - Send via webhook
- receiver

# Notification Pipeline

**Silence**

**Do not  
continue**

**Wait**

**Position in  
cluster  
multiplied  
by 5  
seconds**

**Dedup**

**Has  
notification  
already  
been sent?**

**Send**

**Send  
notification  
via favorite  
provider**

**Gossip**

**Tell other  
peers  
notification  
has been  
sent**

# What is gossiped?

- Yes
  - Sent notifications
  - Silences
- No
  - Received alerts



# How? CRDTs!

- Conflict-free replicated data type
- Associativity ( $a+(b+c)=(a+b)+c$ )
- Commutativity ( $a+b=b+a$ )
- Idempotence ( $a+a=a$ )
- Well suited for AP systems

# How? mesh by Weaveworks

- Eventually consistent
- LWW-element-set
- Mergeable log of records
- Merges based on UID
  - On conflict latest timestamp wins

# Why not etcd?

- Simple operation
  - Less moving pieces
  - Single binary
- Want: AP not CP

# Silences

# Create Silences

Create Silence

**Alertmanager 0**

**Silences  
Database**

ID	Values
1	Query, Start, End
2	Query, Start, End

Gossip Delta  
ID: 2 ...

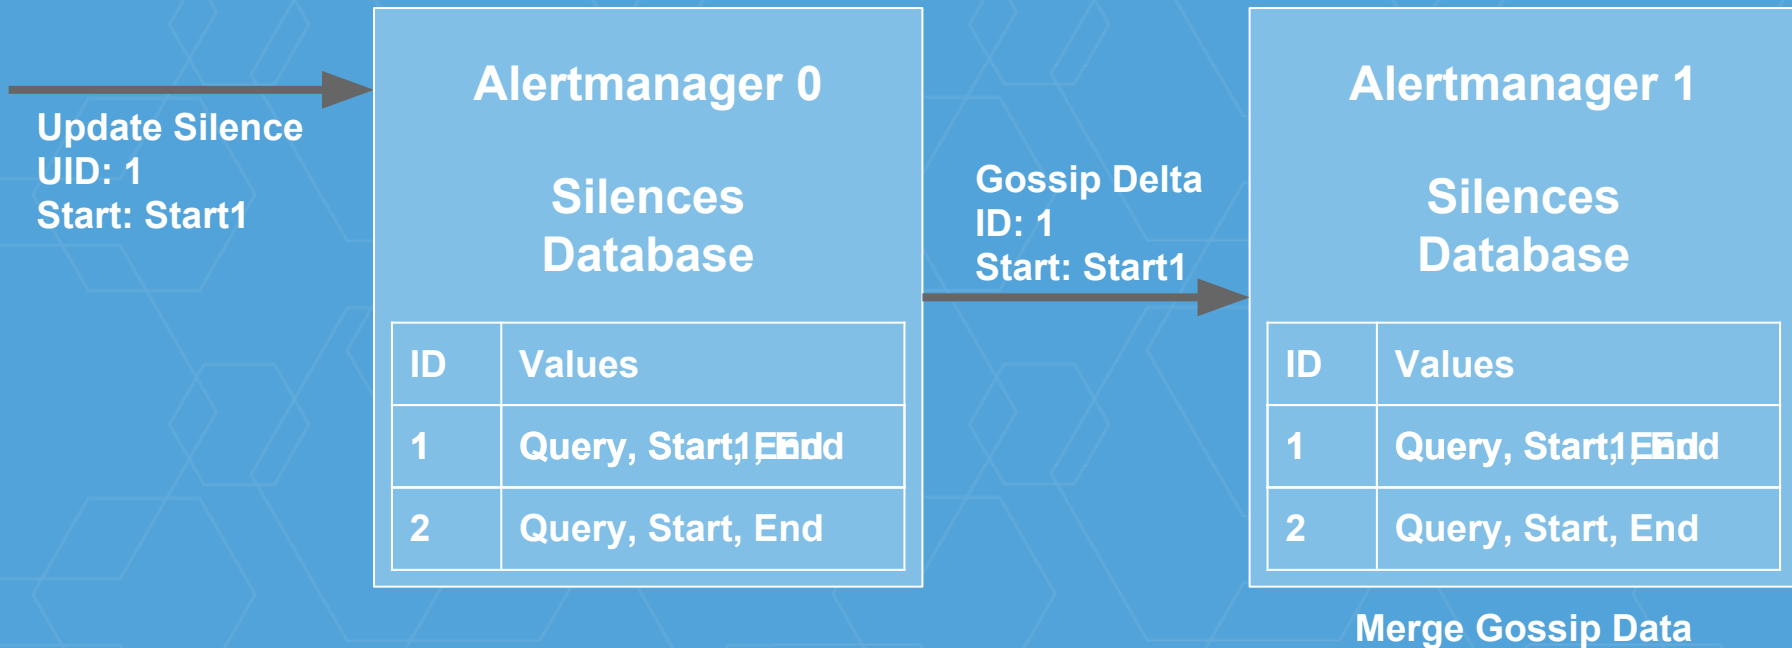
**Alertmanager 1**

**Silences  
Database**

ID	Values
1	Query, Start, End
2	Query, Start, End

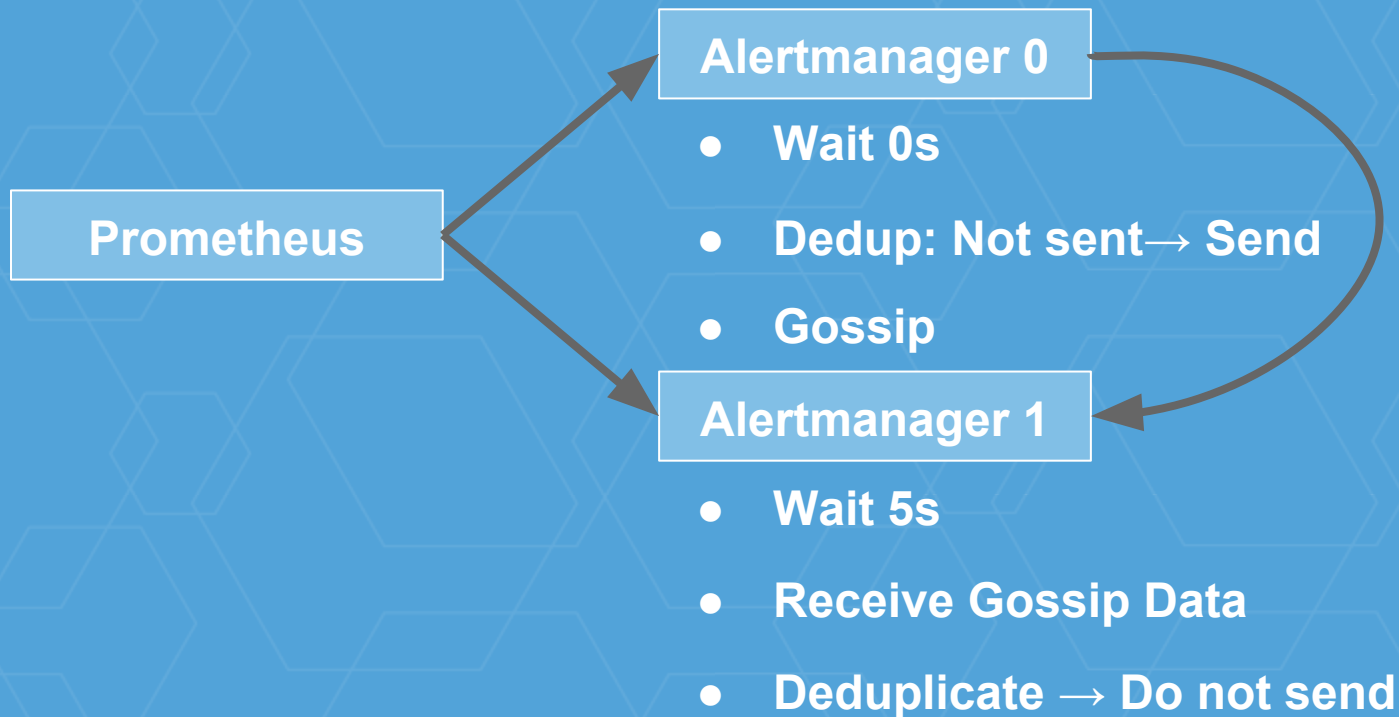
Merge Gossip Data

# Update Silences



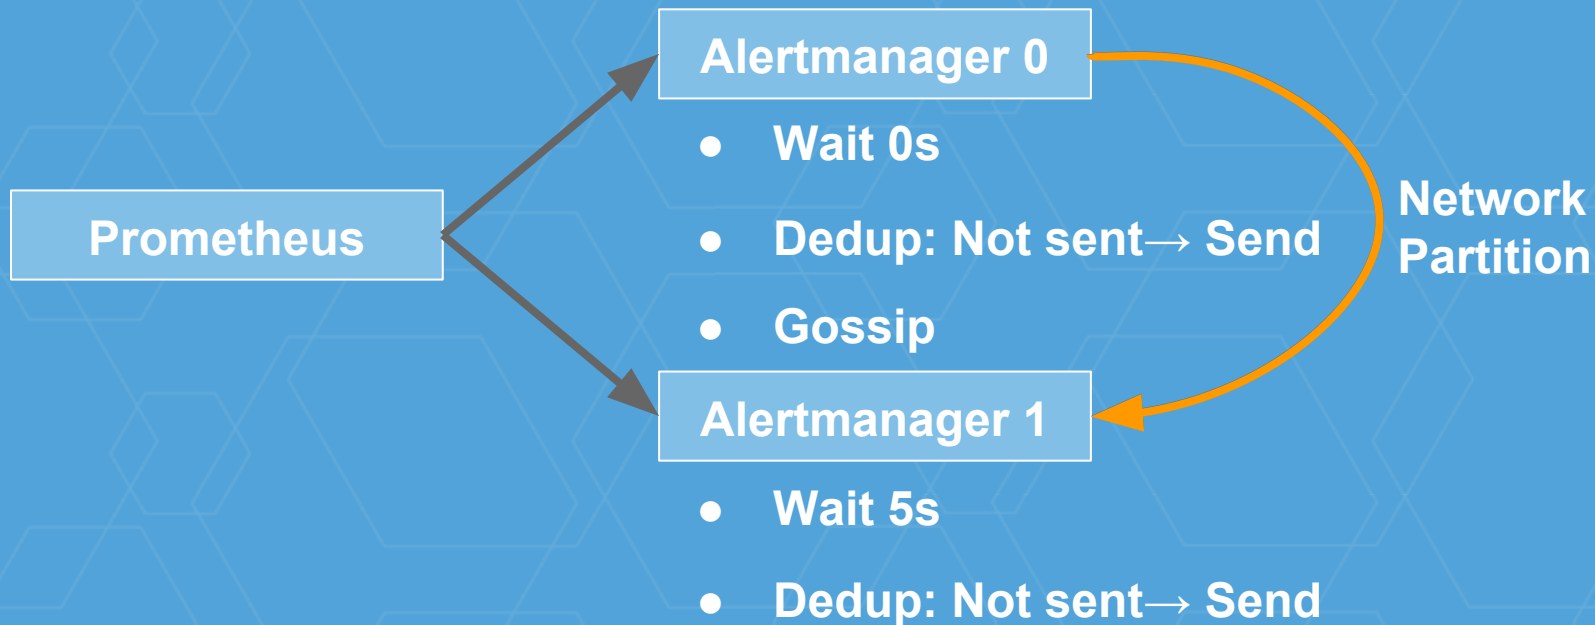
# Notification Log

# Non silenced alert example





# Gossip Partition



# Notification Log

Alert Firing

Alertmanager 0

Notification  
Log

UID	Values
1	IsResolved,TS,...
2	IsResolved,TS,...

Gossip Delta  
UID: 2 ...

Alertmanager 1

Notification  
Log

UID	Values
1	IsResolved,TS,...
2	IsResolved,TS,...

Merge Gossip Data

# Group Key

```
global:
  resolve_timeout: 5m

route:
  group_by: ['job']
  group_wait: 10s
  group_interval: 10s
  repeat_interval: 1h
  receiver: 'webhook'
receivers:
- name: 'webhook'
  webhook_configs:
  - url: 'http://127.0.0.1:5001/'
```

- Group at runtime
  - By Group By labels
- XOR with Route
- Concat with Receiver

The background is a solid blue color with a repeating pattern of white-outlined hexagons of various sizes, creating a honeycomb-like texture.

**DEMO!**

# Thanks!

## QUESTIONS?

[frederic.branczyk@coreos.com](mailto:frederic.branczyk@coreos.com)

GitHub: @brancz

Twitter: @fredbrancz

## LONGER CHAT?

Let's talk!

#prometheus on Freenode

More events: [coreos.com/community](https://coreos.com/community)

# We're hiring: [coreos.com/careers](https://coreos.com/careers)

*ALSO IN BERLIN!*